

Chapter 3: Data Wrangling

Mark Andrews

Psychology Department, Nottingham Trent University

✉ `mark.andrews@ntu.ac.uk`

Data wrangling with the `tidyverse`

- ▶ There are many tools in R for doing data wrangling.
- ▶ Here, we will focus of a core set of inter-related `tidyverse` tools.
- ▶ These include the commands available in the `dplyr` package, particularly its so-called *verbs* such as the following:
 - ▶ `select`
 - ▶ `rename`
 - ▶ `slice`
 - ▶ `filter`
 - ▶ `mutate`
 - ▶ `arrange`
 - ▶ `group_by`
 - ▶ `summarize`

Data wrangling with the tidyverse

- ▶ In addition, `dplyr` provides tools for merging and joining data sets.
- ▶ Next, there are the tools in the `tidyr` package, particularly the following:
 - ▶ `gather`
 - ▶ `spread`
 - ▶ `unite`
 - ▶ `separate`
 - ▶ `pivot_longer`
 - ▶ `pivot_wider`
- ▶ Other packages such as `lubridate` and `stringr` provide essential tools for dealing with and manipulating dates and strings, respectively.
- ▶ All of these tools can be combined together using the `%>%` pipe operator

The dplyr verbs

As an example data set, we will use the data contained in the file `blp-trials-short.txt`.

```
blp <- read_csv('data/blp-trials-short.txt')
```

Select variables with `select`

The `dplyr` command `select` allows us to select columns from a data frame. For example, if we just want `participant`, `lex`, `resp`, and `rt`, then we would do the following.

```
select(blp, participant, lex, resp, rt)
```

```
#> # A tibble: 1,000 x 4
```

```
#>   participant lex    resp    rt
```

```
#>           <dbl> <chr> <chr> <dbl>
```

```
#> 1           20 N      N      977
```

```
#> 2            9 N      N      565
```

```
#> 3           47 N      N      562
```

```
#> 4          103 N      N      572
```

```
#> 5           45 W      W      659
```

```
#> 6           73 W      W      538
```

```
#> 7           24 W      W      626
```

```
#> 8           11 W      W      566
```

```
#> 9           32 W      W      922
```

```
#> 10          96 N      N      555
```

```
#> # ... with 990 more rows
```

Select variables with `select`

We can select a range of variables by specifying the first and last variables in the range with a `:` between them as follows.

```
select(blp, spell:prev.rt)
#> # A tibble: 1,000 x 4
#>   spell      resp      rt prev.rt
#>   <chr>      <chr> <dbl>   <dbl>
#> 1 staud      N      977     511
#> 2 dinbuss    N      565     765
#> 3 snilling   N      562     496
#> 4 gancens    N      572     656
#> 5 filled     W      659     981
#> 6 journals   W      538    1505
#> 7 apache     W      626     546
#> 8 flake      W      566     717
#> 9 reliefs    W      922    1471
#> 10 sarves    N      555     806
#> # ... with 990 more rows
```

Select variables with `select`

We can also select a range of variables using indices as in the following example.

```
select(blp, 2:5) # columns 2 to 5
#> # A tibble: 1,000 x 4
#>   lex      spell      resp      rt
#>   <chr> <chr>      <chr> <dbl>
#> 1 N      staud      N      977
#> 2 N      dinbuss    N      565
#> 3 N      snilling   N      562
#> 4 N      gancens    N      572
#> 5 W      filled     W      659
#> 6 W      journals   W      538
#> 7 W      apache     W      626
#> 8 W      flake      W      566
#> 9 W      reliefs    W      922
#> 10 N     sarves     N      555
#> # ... with 990 more rows
```

Select variables with `select`

We can select variables according to the character or characters that they begin with. For example, we select all variables that begin with `p` as follows.

```
select(blp, starts_with('p'))  
#> # A tibble: 1,000 x 2  
#>   participant prev.rt  
#>   <dbl>     <dbl>  
#> 1         20      511  
#> 2          9      765  
#> 3         47      496  
#> 4        103      656  
#> 5         45      981  
#> 6         73     1505  
#> 7         24      546  
#> 8         11      717  
#> 9         32     1471  
#> 10        96      806  
#> # ... with 990 more rows
```


Select variables with `select`

Or we can select variables by the characters they end with.

```
select(blp, ends_with('t'))  
#> # A tibble: 1,000 x 3  
#>   participant      rt prev.rt  
#>       <dbl> <dbl>   <dbl>  
#> 1         20   977     511  
#> 2          9   565     765  
#> 3         47   562     496  
#> 4        103   572     656  
#> 5         45   659     981  
#> 6         73   538    1505  
#> 7         24   626     546  
#> 8         11   566     717  
#> 9         32   922    1471  
#> 10        96   555     806  
#> # ... with 990 more rows
```

Select variables with `select`

We can select variables that contain a certain set of characters in any position. For example, the following selects variables whose names contain the string `rt`.

```
select(blp, contains('rt'))  
#> # A tibble: 1,000 x 4  
#>   participant      rt prev.rt rt.raw  
#>   <dbl> <dbl>   <dbl> <dbl>  
#> 1         20    977     511    977  
#> 2          9    565     765    565  
#> 3         47    562     496    562  
#> 4        103    572     656    572  
#> 5         45    659     981    659  
#> 6         73    538    1505    538  
#> 7         24    626     546    626  
#> 8         11    566     717    566  
#> 9         32    922    1471    922  
#> 10        96    555     806    555  
#> # ... with 990 more rows
```

Select variables with `select`

We can also match by regular expressions. For example, the regular expression `^rt|rt$` will match the `rt` if it begins or ends a string. Therefore, we can select the variables that contain `rt`, where the string `rt` means reaction time, as follows.

```
select(blp, matches('^rt|rt$'))
```

```
#> # A tibble: 1,000 x 3
```

```
#>       rt prev.rt rt.raw
```

```
#>   <dbl>   <dbl> <dbl>
```

```
#> 1    977     511    977
```

```
#> 2    565     765    565
```

```
#> 3    562     496    562
```

```
#> 4    572     656    572
```

```
#> 5    659     981    659
```

```
#> 6    538    1505    538
```

```
#> 7    626     546    626
```

```
#> 8    566     717    566
```

```
#> 9    922    1471    922
```

```
#> 10   555     806    555
```

```
#> # ... with 990 more rows
```

Remove variables with `select`

We can use `select` to *remove* variables as well as select them. To remove a variable, we precede its name with a minus sign.

```
select(blp, -participant) # remove `participant`  
#> # A tibble: 1,000 x 6  
#>   lex      spell      resp      rt prev.rt rt.raw  
#>   <chr> <chr>      <chr> <dbl>   <dbl>   <dbl>  
#> 1 N      staud      N      977     511     977  
#> 2 N      dinbuss    N      565     765     565  
#> 3 N      snilling  N      562     496     562  
#> 4 N      gancens    N      572     656     572  
#> 5 W      filled     W      659     981     659  
#> 6 W      journals  W      538    1505     538  
#> 7 W      apache     W      626     546     626  
#> 8 W      flake      W      566     717     566  
#> 9 W      reliefs    W      922    1471     922  
#> 10 N     sarves     N      555     806     555  
#> # ... with 990 more rows
```

Remove variables with `select`

Just as we selected ranges or sets of variables above, we can remove them by preceding their selection functions with minus signs.

```
select(blp, -(2:6))  
#> # A tibble: 1,000 x 2  
#>   participant rt.raw  
#>   <dbl> <dbl>  
#> 1      20    977  
#> 2       9    565  
#> 3      47    562  
#> 4     103    572  
#> 5      45    659  
#> 6      73    538  
#> 7      24    626  
#> 8      11    566  
#> 9      32    922  
#> 10     96    555  
#> # ... with 990 more rows
```

Remove variables with `select`

Or, as another example, we can remove the variables that contain the string `rt` as follows.

```
select(blp, -contains('rt'))  
#> # A tibble: 1,000 x 3  
#>   lex      spell      resp  
#>   <chr> <chr>      <chr>  
#> 1 N      staud      N  
#> 2 N      dinbuss    N  
#> 3 N      snilling   N  
#> 4 N      gancens    N  
#> 5 W      filled     W  
#> 6 W      journals   W  
#> 7 W      apache     W  
#> 8 W      flake      W  
#> 9 W      reliefs    W  
#> 10 N     sarves     N  
#> # ... with 990 more rows
```

Renaming variables with `select`

When we select individual variables with `select`, we can rename them too, as in the following example.

```
select(blp, subject=participant, reaction_time=rt)
#> # A tibble: 1,000 x 2
#>   subject reaction_time
#>   <dbl>         <dbl>
#> 1         20         977
#> 2          9         565
#> 3         47         562
#> 4        103         572
#> 5         45         659
#> 6         73         538
#> 7         24         626
#> 8         11         566
#> 9         32         922
#> 10        96         555
#> # ... with 990 more rows
```

Renaming variables with `rename`

If we want to rename some variables, and get a data frame with all variables, including the renamed ones, we should use `rename`.

```
rename(blp, subject=participant, reaction_time=rt)
#> # A tibble: 1,000 x 7
#>   subject lex spell resp reaction_time prev.rt rt.raw
#>   <dbl> <chr> <chr> <chr>          <dbl>   <dbl> <dbl>
#> 1      20 N   staud   N           977     511   977
#> 2       9 N   dinbuss N           565     765   565
#> 3      47 N   snilling N           562     496   562
#> 4     103 N   gancens N           572     656   572
#> 5      45 W   filled   W           659     981   659
#> 6      73 W   journals W           538    1505   538
#> 7      24 W   apache   W           626     546   626
#> 8      11 W   flake    W           566     717   566
#> 9      32 W   reliefs W           922    1471   922
#> 10     96 N   sarves   N           555     806   555
#> # ... with 990 more rows
```


Selecting observations by indices with `slice`

We use `slice` to select observations by their indices. For example, to select rows 10, 20, 50, 100, 500, we would simply do the following.

```
slice(blp, c(10, 20, 50, 100, 500))
```

```
#> # A tibble: 5 x 7
```

#>	participant	lex	spell	resp	rt	prev.rt	rt.raw
#>		<dbl>	<chr>	<chr>	<chr>	<dbl>	<dbl>
#> 1	96	N	sarves	N	555	806	555
#> 2	46	W	mirage	W	778	571	778
#> 3	72	N	gright	N	430	675	430
#> 4	3	W	gleam	W	361	370	361
#> 5	92	W	coaxes	W	699	990	699

Selecting observations by indices with `slice`

Given that, for example, `10:100` would list the integers 10 to 100 inclusive, we can select just these observations as follows.

```
slice(blp, 10:100)
```

```
#> # A tibble: 91 x 7
```

#>	participant	lex	spell	resp	rt	prev.rt	rt.raw
#>		<dbl>	<chr>	<chr>	<chr>	<dbl>	<dbl>
#> 1	96	N	sarves	N	555	806	555
#> 2	82	W	deceits	W	657	728	657
#> 3	37	W	nothings	N	NA	552	712
#> 4	52	N	chuespies	N	427	539	427
#> 5	96	N	mowny	N	1352	1020	1352
#> 6	96	N	cranned	N	907	573	907
#> 7	89	N	flud	N	742	834	742
#> 8	3	N	bromble	N	523	502	523
#> 9	7	N	trubbles	N	782	458	782
#> 10	35	N	playfound	N	643	663	643

```
#> # ... with 81 more rows
```

De-selecting observations by indices with `slice`

Just as we did with `select`, we can precede the indices with a minus sign to drop the corresponding observations. Thus, for example, we can drop the first 10 observations as follows.

```
slice(blp, -(1:10))
```

#> # A tibble: 990 x 7

#>	participant	lex	spell	resp	rt	prev.rt	rt.raw
#>		<dbl>	<chr>	<chr>	<chr>	<dbl>	<dbl>
#> 1	82	W	deceits	W	657	728	657
#> 2	37	W	nothings	N	NA	552	712
#> 3	52	N	chuespies	N	427	539	427
#> 4	96	N	mowny	N	1352	1020	1352
#> 5	96	N	cranned	N	907	573	907
#> 6	89	N	flud	N	742	834	742
#> 7	3	N	bromble	N	523	502	523
#> 8	7	N	trubbles	N	782	458	782
#> 9	35	N	playfound	N	643	663	643
#> 10	46	W	mirage	W	778	571	778

#> # ... with 980 more rows

Selecting observations by condition with `filter`

The `filter` command is a powerful means to filter observations according to their values. For example, we can select all the observations where the `lex` variable is N as follows.

```
filter(blp, lex == 'N')
```

```
#> # A tibble: 502 x 7
```

```
#>   participant lex    spell    resp    rt prev.rt rt.raw
#>   <dbl> <chr> <chr>    <chr> <dbl>    <dbl> <dbl>
#> 1         20 N    staud    N      977      511    977
#> 2          9 N    dinbuss  N      565      765    565
#> 3         47 N    snilling N      562      496    562
#> 4        103 N    gancens  N      572      656    572
#> 5         96 N    sarves   N      555      806    555
#> 6         52 N    chuespies N      427      539    427
#> 7         96 N    mowny    N     1352     1020   1352
#> 8         96 N    cranned  N      907      573    907
#> 9         89 N    flud     N      742      834    742
#> 10         3 N    bromble  N      523      502    523
#> # ... with 492 more rows
```

Selecting observations by condition with `filter`

We can also filter by multiple conditions by listing each one with commas between them. For example, the following gives us the observations where `lex` has the value of N and `resp` has the value of W.

```
filter(blp, lex == 'N', resp=='W')  
#> # A tibble: 35 x 7  
#>   participant lex    spell    resp    rt prev.rt rt.raw  
#>   <dbl> <chr> <chr>    <chr> <dbl>    <dbl> <dbl>  
#> 1         73 N    bunding W        NA     978    1279  
#> 2         63 N    gallays W        NA     589     923  
#> 3         50 N    droper  W        NA     741     573  
#> 4          6 N    flooder W        NA     524     557  
#> 5         73 N    khantum W        NA     623    1355  
#> 6         81 N    seaped  W        NA     765     691  
#> 7         43 N    gafers  W        NA     556     812  
#> 8        101 N    winchers W        NA     632     852  
#> 9         81 N    flaged  W        NA     674     609  
#> 10        11 N    frocker W        NA     653     665  
#> # ... with 25 more rows
```

Selecting observations by condition with `filter`

The following gives us those observations where `lex` has the value of N and `resp` has the value of W and `rt.raw` is less than or equal to 500.

```
filter(blp, lex == 'N', resp=='W', rt.raw <= 500)
#> # A tibble: 5 x 7
#>   participant lex    spell    resp    rt prev.rt rt.raw
#>   <dbl> <chr> <chr>    <chr> <dbl>    <dbl> <dbl>
#> 1      28 N    cown      W      NA     680    498
#> 2      17 N    beeched  W      NA     450    469
#> 3      29 N    conform W      NA     495    497
#> 4      35 N    blear   W      NA     592    461
#> 5      89 N    stumming W      NA     571    442
```

Selecting observations by condition with `filter`

The previous command is equivalent to making a conjunction of conditions using `&` as follows.

```
filter(blp, lex == 'N' & resp=='W' & rt.raw <= 500)
```

```
#> # A tibble: 5 x 7
```

```
#>   participant lex      spell      resp      rt prev.rt rt.raw  
#>           <dbl> <chr> <chr>      <chr> <dbl>    <dbl> <dbl>  
#> 1          28 N      cown      W      NA      680    498  
#> 2          17 N    beeched    W      NA      450    469  
#> 3          29 N    conform    W      NA      495    497  
#> 4          35 N    blear     W      NA      592    461  
#> 5          89 N    stumming  W      NA      571    442
```

Selecting observations by condition with `filter`

We can make a *disjunction* of conditions for filtering using the logical-or symbol `|`. For example, to filter observation where the `rt.raw` was either less than 500 or greater than 1000, we can do the following.

```
filter(blp, rt.raw < 500 | rt.raw > 1000)
```

```
#> # A tibble: 296 x 7
```

```
#>   participant lex    spell    resp    rt prev.rt rt.raw
#>   <dbl> <chr> <chr>    <chr> <dbl>   <dbl>   <dbl>
#> 1      52 N    chuespies N      427     539     427
#> 2      96 N    mowny     N     1352    1020    1352
#> 3      28 W    stelae    N      NA     678     497
#> 4      85 W    forewarned N      NA     525     350
#> 5      24 W    owl      W     470     535     470
#> 6      97 W    soda       W     436     447     436
#> 7      81 N    fugate     N     425     403     425
#> 8     105 N    pamps      N      NA     884    1494
#> 9      27 W    outgrowth  N      NA     633    1014
#> 10     82 W    kitty      W     431     476     431
#> # ... with 286 more rows
```


Changing variables and values with mutate

To create a new variable `is_accurate` that takes the value of `TRUE` whenever `lex` and `resp` have the same value, we can do the following:

```
mutate(blpl, acc = lex == resp)
```

```
#> # A tibble: 1,000 x 8
```

```
#>   participant lex      spell      resp      rt prev.rt rt.raw acc
#>   <dbl> <chr> <chr>      <chr> <dbl>    <dbl> <dbl> <lgl>
#> 1      20 N      staud      N      977      511    977 TRUE
#> 2       9 N    dinbuss      N      565      765    565 TRUE
#> 3      47 N    snilling      N      562      496    562 TRUE
#> 4     103 N    gancens      N      572      656    572 TRUE
#> 5      45 W    filled      W      659      981    659 TRUE
#> 6      73 W    journals      W      538     1505    538 TRUE
#> 7      24 W    apache      W      626      546    626 TRUE
#> 8      11 W    flake      W      566      717    566 TRUE
#> 9      32 W    reliefs      W      922     1471    922 TRUE
#> 10     96 N    sarves      N      555      806    555 TRUE
#> # ... with 990 more rows
```

Changing variables and values with mutate

As another example, we can create a new variable that gives the length of the word given by the `spell` variable.

```
mutate(blp, len = str_length(spell))
```

```
#> # A tibble: 1,000 x 8
```

```
#>   participant lex    spell    resp    rt prev.rt rt.raw    le
#>   <dbl> <chr> <chr>    <chr> <dbl>    <dbl> <dbl> <int>
```

```
#> 1         20 N    staud    N    977    511    977
```

```
#> 2          9 N   dinbuss    N    565    765    565
```

```
#> 3         47 N   snilling    N    562    496    562
```

```
#> 4        103 N   gancens    N    572    656    572
```

```
#> 5         45 W   filled    W    659    981    659
```

```
#> 6         73 W   journals    W    538   1505    538
```

```
#> 7         24 W   apache    W    626    546    626
```

```
#> 8         11 W   flake    W    566    717    566
```

```
#> 9         32 W   reliefs    W    922   1471    922
```

```
#> 10        96 N   sarves    N    555    806    555
```

```
#> # ... with 990 more rows
```

Changing variables and values with mutate

We can also create multiple new variable at the same time as in the following example.

```
mutate(blp,
      acc = lex == resp,
      fast = rt.raw < mean(rt.raw, na.rm=TRUE))
#> # A tibble: 1,000 x 9
#>   participant lex    spell    resp    rt prev.rt rt.raw acc
#>   <dbl> <chr> <chr>    <chr> <dbl>    <dbl> <dbl> <lgl>
#> 1         20 N    staud    N     977      511    977 TRUE
#> 2          9 N   dinbuss  N     565      765    565 TRUE
#> 3         47 N   snilling N     562      496    562 TRUE
#> 4        103 N   gancens  N     572      656    572 TRUE
#> 5         45 W   filled   W     659      981    659 TRUE
#> 6         73 W   journals W     538     1505    538 TRUE
#> 7         24 W   apache   W     626      546    626 TRUE
#> 8         11 W   flake    W     566      717    566 TRUE
#> 9         32 W   reliefs  W     922     1471    922 TRUE
#> 10        96 N   sarves   N     555      806    555 TRUE
#> # ... with 990 more rows
```

Sorting observations with arrange

Sorting observations in a data frame is easily accomplished with `arrange`. For example to sort by `participant` and then by `spell`, we would do the following.

```
arrange(blp, participant, spell)
```

```
#> # A tibble: 1,000 x 7
```

```
#>   participant lex    spell    resp    rt prev.rt rt.raw
#>   <dbl> <chr> <chr>    <chr> <dbl>    <dbl> <dbl>
#> 1         1 W    abyss    W      629      683      629
#> 2         1 N    baises    N      524      574      524
#> 3         1 W    carport    W      779      605      779
#> 4         1 N    cellies    N      792      652      792
#> 5         1 W    chafing    W      601      720      601
#> 6         1 N    dametails N      694      635      694
#> 7         1 N    footer    N      789      566      789
#> 8         1 W    gantries    W      644      581      644
#> 9         1 N    hogtush    N      679      568      679
#> 10        1 N    lisedess    N      679      619      679
#> # ... with 990 more rows
```

Sorting observations with arrange

We can sort by the reverse order of any variable by using the `desc` command on the variable. In the following example, we sort by `participant`, and then by `spell` in reverse order.

```
arrange(blp, participant, desc(spell))
```

```
#> # A tibble: 1,000 x 7
```

```
#>   participant lex    spell    resp    rt prev.rt rt.raw
#>   <dbl> <chr> <chr>    <chr> <dbl>    <dbl> <dbl>
#> 1         1 N    wintes    N     545     629     545
#> 2         1 N    treeps    N     607     610     607
#> 3         1 W    squashes W     494     491     494
#> 4         1 N    sinkhicks N     536     519     536
#> 5         1 W    shafting W     553     571     553
#> 6         1 W    month    W     500     498     500
#> 7         1 N    lisedess N     679     619     679
#> 8         1 N    hogtush  N     679     568     679
#> 9         1 W    gantries W     644     581     644
#> 10        1 N    footer  N     789     566     789
#> # ... with 990 more rows
```

Reducing data with `summary`

The `dplyr` package has a function `summarize` (or, equivalently, `summarise`) that applies summarizing functions to variables.

For example, we may calculate some summary statistics of the particular variables as in the following example.

```
summarize(blp,  
           mean_rt = mean(rt, na.rm = T),  
           median_rt = median(rt, na.rm = T),  
           sd_rt.raw = sd(rt.raw, na.rm = T)  
)  
#> # A tibble: 1 x 3  
#>   mean_rt median_rt sd_rt.raw  
#>   <dbl>     <dbl>     <dbl>  
#> 1    638.       588       474.
```

(Note that here it is necessary to use `na.rm = T` to remove the NA values in the variables.)

Reducing data with `summary` and `group_by`

The `summarize` command, and its variants, become considerably more powerful when combined with the `group_by` command. Effectively, `group_by` groups the observations within a data frame according to the values of specified variables. For example, the following command groups `blp` into groups of observations according to value of the `lex` variable.

```
blp_by_lex <- group_by(blp, lex)
```

Reducing data with `summary` and `group_by`

If we now apply `summarize` to this grouped data frame, we will obtain summary statistics for each group, as in the following example.

```
summarize(blp_by_lex, mean = mean(rt, na.rm=T))  
#> # A tibble: 2 x 2  
#>   lex      mean  
#>   <chr> <dbl>  
#> 1 N      638.  
#> 2 W      637.
```


The %>% operator

The %>% operator in R is known as the *pipe*. It is available from the `magrittr` package, which is part of the `tidyverse`. In RStudio, the keyboard shortcut Ctrl+Shift+M types %>%.

To understand pipes, let us begin with a very simple example. The following `primes` variable is a vector of the first 10 prime numbers.

```
primes <- c(2, 3, 5, 7, 11, 13, 17, 19, 23, 29)
```

We can calculate the sum of `primes` as follows.

```
sum(primes)
#> [1] 129
```

We may then calculate the square root of this sum.

```
sqrt(sum(primes))
#> [1] 11.35782
```

We may then calculate the logarithm of this square root.

```
log(sqrt(sum(primes)))
#> [1] 2.429906
```

The %>% operator

The %>% is *syntactic sugar* that reexpresses nested functions as sequences.

Returning to some of our examples above, we will see how they can be rewritten with pipes. In each case, we will precede the piped version with a comment showing its original version.

```
# sum(primes)
primes %>% sum()
#> [1] 129
```

```
# sum(primes, na.rm=T)
primes %>% sum(na.rm=T)
#> [1] 129
```

```
# log(sqrt(sum(primes)))
primes %>% sum() %>% sqrt() %>% log()
#> [1] 2.429906
```

Reshaping with `pivot_longer` and `pivot_wider`

A so-called *tidy* data set is a data set where all rows are observations, all columns are variables, and each variable is a single value.

Consider the following data frame.

```
recall_df <- read_csv('data/repeated_measured_a.csv')
```

```
recall_df
```

```
#> # A tibble: 5 x 4
```

```
#>   Subject    Neg    Neu    Pos
```

```
#>   <chr>    <dbl> <dbl> <dbl>
```

```
#> 1 Faye      26     12     42
```

```
#> 2 Jason      29      8     35
```

```
#> 3 Jim       32     15     45
```

```
#> 4 Ron       22     10     38
```

```
#> 5 Victor    30     13     40
```

In this data frame, for each subject, we have three values, which are their scores on a memory test in three different conditions of an experiment. The conditions are `Neg` (negative), `Neu` (neutral), `Pos` (positive). However, each column is not a variable. The `Neg`, `Neu`, `Pos` are, in fact, *values* of a variable, namely the condition of the experiment.

Reshaping with `pivot_longer`

To tidy this data frame, we need a variable for the subject, another for the experiment's condition, and another for the memory score for the corresponding subject in the corresponding condition. To do so, we perform what is sometimes known as a *wide to long* transformation. The `tidyr` package has a function `pivot_longer` for this transformation.

To use `pivot_longer`, we must specify the variables (using the `cols` argument) that we want to pivot from wide to long. Next, we must provide a name for the column that will indicate the experimental condition. Finally, we must provide a name for the column that will indicate the memory scores.

Reshaping with pivot_longer

Here is the necessary code:

```
recall_long <- pivot_longer(  
  recall_df, cols = -Subject, names_to = 'condition',  
  values_to = 'score')
```

```
recall_long
```

```
#> # A tibble: 15 x 3
```

```
#>   Subject condition score
```

```
#>   <chr>    <chr>    <dbl>
```

```
#> 1 Faye     Neg      26
```

```
#> 2 Faye     Neu      12
```

```
#> 3 Faye     Pos      42
```

```
#> 4 Jason    Neg      29
```

```
#> 5 Jason    Neu       8
```

```
#> 6 Jason    Pos      35
```

```
#> 7 Jim      Neg      32
```

```
#> 8 Jim      Neu      15
```

```
#> 9 Jim      Pos      45
```

```
#> 10 Ron     Neg      22
```

```
#> 11 Ron     Neu      10
```

```
#> 12 Ron     Pos      38
```

Reshaping with `pivot_wider`

The inverse of a `pivot_longer` is a `pivot_wider`. It is very similar to `pivot_longer` and we use `names_from` and `values_from` in the opposite sense to `names_to` and `values_to`.

```
pivot_wider(recall_long,
             names_from = 'condition', values_from = 'score')
```

#> # A tibble: 5 x 4

<i>#></i>	<i>Subject</i>	<i>Neg</i>	<i>Neu</i>	<i>Pos</i>
<i>#></i>	<i><chr></i>	<i><dbl></i>	<i><dbl></i>	<i><dbl></i>
<i>#> 1</i>	<i>Faye</i>	<i>26</i>	<i>12</i>	<i>42</i>
<i>#> 2</i>	<i>Jason</i>	<i>29</i>	<i>8</i>	<i>35</i>
<i>#> 3</i>	<i>Jim</i>	<i>32</i>	<i>15</i>	<i>45</i>
<i>#> 4</i>	<i>Ron</i>	<i>22</i>	<i>10</i>	<i>38</i>
<i>#> 5</i>	<i>Victor</i>	<i>30</i>	<i>13</i>	<i>40</i>